

Doubly Compressed Sparse Column (DCSC) Storage

Kalyani Gadgil

December 6, 2017

Need for Hypersparse Matrices

Matrix is hypersparse if number of non-zero elements is much less than the dimensions of the matrix

$$nnz < n$$

Hypersparse matrices arise after 2-dimensional block data decomposition of matrices for parallel processing like in SUMMA.

Storage Complexity of Sparse Matrices for SUMMA

- 1 Blocks of size $(n/\sqrt{p}) \times (n/\sqrt{p})$
where n - matrix dim and p - number of processors
- 2 Storing each of these submatrices in CSC format
 $O(n\sqrt{p} + nnz)$
- 3 Storing whole matrix $O(n + nnz)$ on single processor
- 4 Storing matrix in DCSC format requires $O(nnz)$

Arrays in DCSC

- 1 JC in CSC allows fast access to columns but not rows
- 2 solution could be to store CSR as well but that doubles storage
- 3 information theoretic solution is to remove unnecessary repetitions from JC \Rightarrow CP array formed
- 4 CP array contains pointers to row indices of nnz elements
- 5 compressing JC array from $n+1$ to nzc (number of columns containing atleast one non-zero element), leads to indexing issues
- 6 this \Rightarrow form an auxilliary array (AUX) to store pointers to nonzero columns

Access element at $A(i,j)$

Pointers mean indices or location of value inside an array/vector. Not pointers to physical memory addresses of values.

- 1 find out which column chunk the element belongs to using chunk sizes determined as $\lceil cf \rceil = (n + 1)/nzc$
- 2 get index $AUX[j/chunk].. + 1$ which returns subarray of nonzero columns in that chunk
- 3 Search this subarray of JC for j ; if found, store the index pos
- 4 if found, we know right now that there is some nnz element in this row. Now we search for the specific row we need
- 5 Use $CP[pos].. + 1$ to get subarray of all elements in a particular row
- 6 search subarray of IR for i ; if found, store index as $posc$
- 7 Use index to get value at that position $NUM[posc]$

DCSC Storage

Matrix A

$$\begin{matrix} & 0 & 1 & 2 & 3 & 4 & 5 \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 4 \end{pmatrix} \end{matrix}$$

$$n = 6, \text{nnz} = 4$$

$$\lceil cf \rceil = n + 1/nzc = \lceil 1.25 \rceil = 2$$

$$\text{NUM} = [1, 2, 3, 4]$$

$$\text{row idx (IR)} = [0, 3, 2, 5]$$

$$\text{JC from CSC} = [0, 0, | 2, 3, | 3, 3, | 4]$$

$$\text{col idx (JC)} = [1, 1, 2, 5]$$

DCSC Storage Walkthrough

Matrix A

	0	1	2	3	4	5
0	0	1	0	0	0	0
1	0	0	0	0	0	0
2	0	0	3	0	0	0
3	0	2	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	4

$$\text{NUM} = [1, 2]$$

$$\text{IR} = [0, 3]$$

$$\text{idx_IR} = [0, 1]$$

CP stores ptrs to idx of IR when
col changes

$$\text{CP} = [0]$$

JC stores column indices

$$\text{JC} = [1]$$

$$\text{idx_JC} = [0, 1]$$

AUX stores one ptr to idx of JC
for each chunk

$$\text{AUX} = [0]$$

DCSC (contd. 1)

Matrix A

$$\begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 4 \end{pmatrix} \end{matrix}$$

$$\text{NUM} = [1, 2, 3]$$

$$\text{IR} = [0, 3, 2]$$

$$\text{idx_IR} = [0, 1, 2]$$

CP stores ptrs to idx of IR when
col changes

$$\text{CP} = [0, 2]$$

JC stores column indices

$$\text{JC} = [1, 2]$$

$$\text{idx_JC} = [0, 1, 2]$$

AUX stores one ptr to idx of JC
for each chunk

$$\text{AUX} = [0]$$

DCSC (contd. 2)

		Matrix A				
	0	1	2	3	4	5
0	0	1	0	0	0	0
1	0	0	0	0	0	0
2	0	0	3	0	0	0
3	0	2	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	4

$$\text{NUM} = [1, 2, 3, \boxed{4}]$$

$$\text{IR} = [0, 3, 2, \boxed{5}]$$

$$\text{idx_IR} = [0, 1, 2, \boxed{3}]$$

CP stores ptrs to idx of IR when
col changes

$$\text{CP} = [0, 2, \boxed{3}]$$

JC stores column indices

$$\text{JC} = [1, 2, \boxed{5}]$$

$$\text{idx_JC} = [0, 1, \boxed{2}, 3]$$

AUX stores one ptr to idx of JC
for each chunk

$$\text{AUX} = [0, \boxed{2}]$$

DCSC (contd. 3)

$$\begin{array}{c} \text{Matrix A} \\ \begin{array}{cccccc} & 0 & 1 & 2 & 3 & 4 & 5 \\ \begin{array}{c} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{array} & \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 4 \end{pmatrix} \end{array} \end{array}$$

End of matrix reached therefore,
nnz added to CP and AUX

$$\text{NUM} = [1, 2, 3, 4]$$

$$\text{IR} = [0, 3, 2, 5]$$

$$\text{idx_IR} = [0, 1, 2, \boxed{3}]$$

CP stores ptrs to idx of IR when
col changes

$$\text{CP} = [0, 2, 3, \boxed{3}]$$

JC stores column indices

$$\text{JC} = [1, 2, 5]$$

$$\text{idx_JC} = [0, 1, 2, \boxed{3}]$$

AUX stores one ptr to idx of JC
for each chunk

$$\text{AUX} = [0, 2, \boxed{3}]$$

CSC example - Column Major

① $O(n + nnz)$ comes from JC array of size $n + 1$

⁰[1] Buluc, A., & Gilbert, J. R. (2008). On the Representation and Multiplication of Hypersparse Matrices.
<https://doi.org/10.1109/IPDPS.2008.4536313>

References

- [1] Buluc, A., & Gilbert, J. R. (2008). On the Representation and Multiplication of Hypersparse Matrices.
<https://doi.org/10.1109/IPDPS.2008.4536313>